

5118020-03 Operating Systems

Homework 2. smalloc

Shin Hong

Overview

2

- Complete `smalloc.c`, an in-house heap memory allocator for C programs
- Point of study
 - memory management API: `mmap()`
- Timelines
 - Apr 30: First announcement & team arrangement
 - May 6-7: Help desks
 - May 13: Second test case releases
 - May 15, 9 PM: Artifact submission deadline (source code)
 - May 16, 9 PM: Presentation submission deadline (video record)

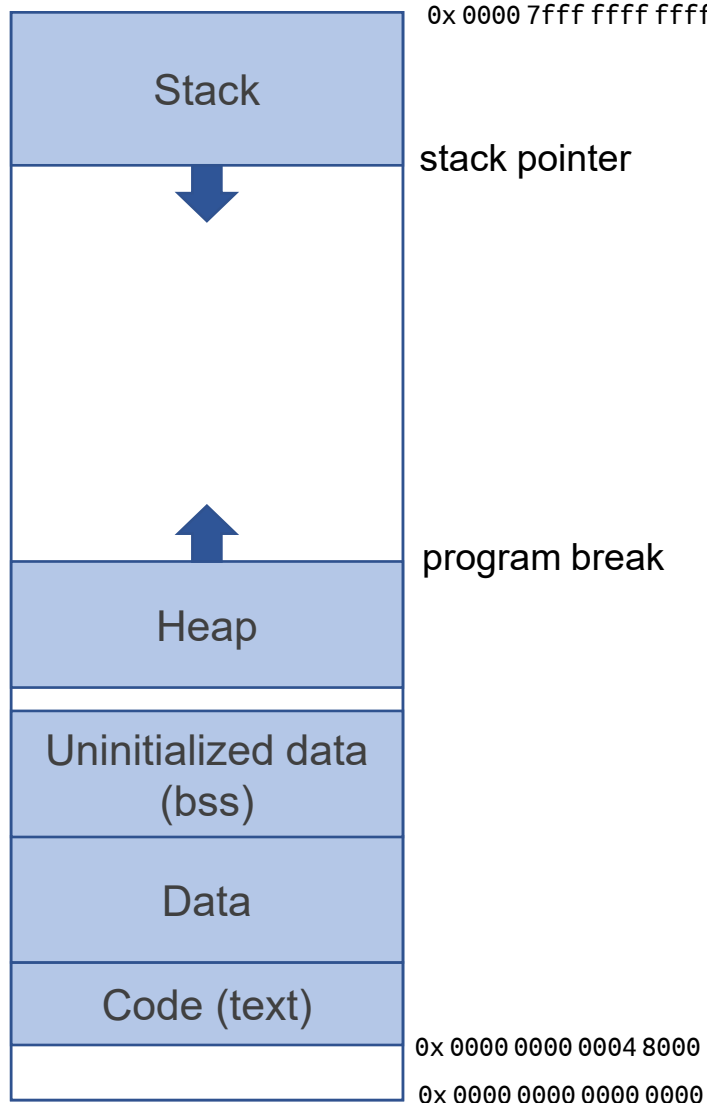
Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02

Background: Segmentation Layout (Linux, x86-64)

3



- `&etext` points to the first address past the end of the text segment
- `&edata` points to the first address past the end of initialized data segment
- `&end` points to the first address past the end of the uninitialized data segment
- `sbrk(0)` returns the first address past the end of the currently given heap segment
- `sbrk(s)` retains additional `s` bytes in heap and returns the starting address.
 - returns null when OS denies the request
- `getpagesize()` returns the number of bytes in a page
- c.f. https://en.wikipedia.org/wiki/X86-64#Virtual_address_space_details

Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02

Background: `mmap()`

- `mmap()` adds one or multiple pages to the current process
 - the main use case is to set up a memory area for memory mapped I/O
 - `mmap()` with option `MAP_ANON` acquires a new pages while not connecting them with any file
 - the size of memory request must be a multiple of page size
 - `getpagesize()`: <https://man7.org/linux/man-pages/man2/getpagesize.2.html>
- `mmap()` manpage: <https://man7.org/linux/man-pages/man2/mmap.2.html>

Homework 2.
smalloc


5118020-03
Operating Systems

2024-05-02

Memory Allocation Management (1/2)

5

- Use `mmap()` to acquire one or more pages as heap space
- Manage the information about the heap space as a linked list
 - a **data region** is a series of consecutive bytes to store user's data
 - a **data header** records the attributes of the corresponding data region, and it is located right before the data region in the memory address
 - the length of the data region
 - whether the data region is “used” or “unused”
 - used: allocated (given to the user) and not yet freed
 - unused: never used or freed
 - the pointer to the next data header (i.e., form a linked list)



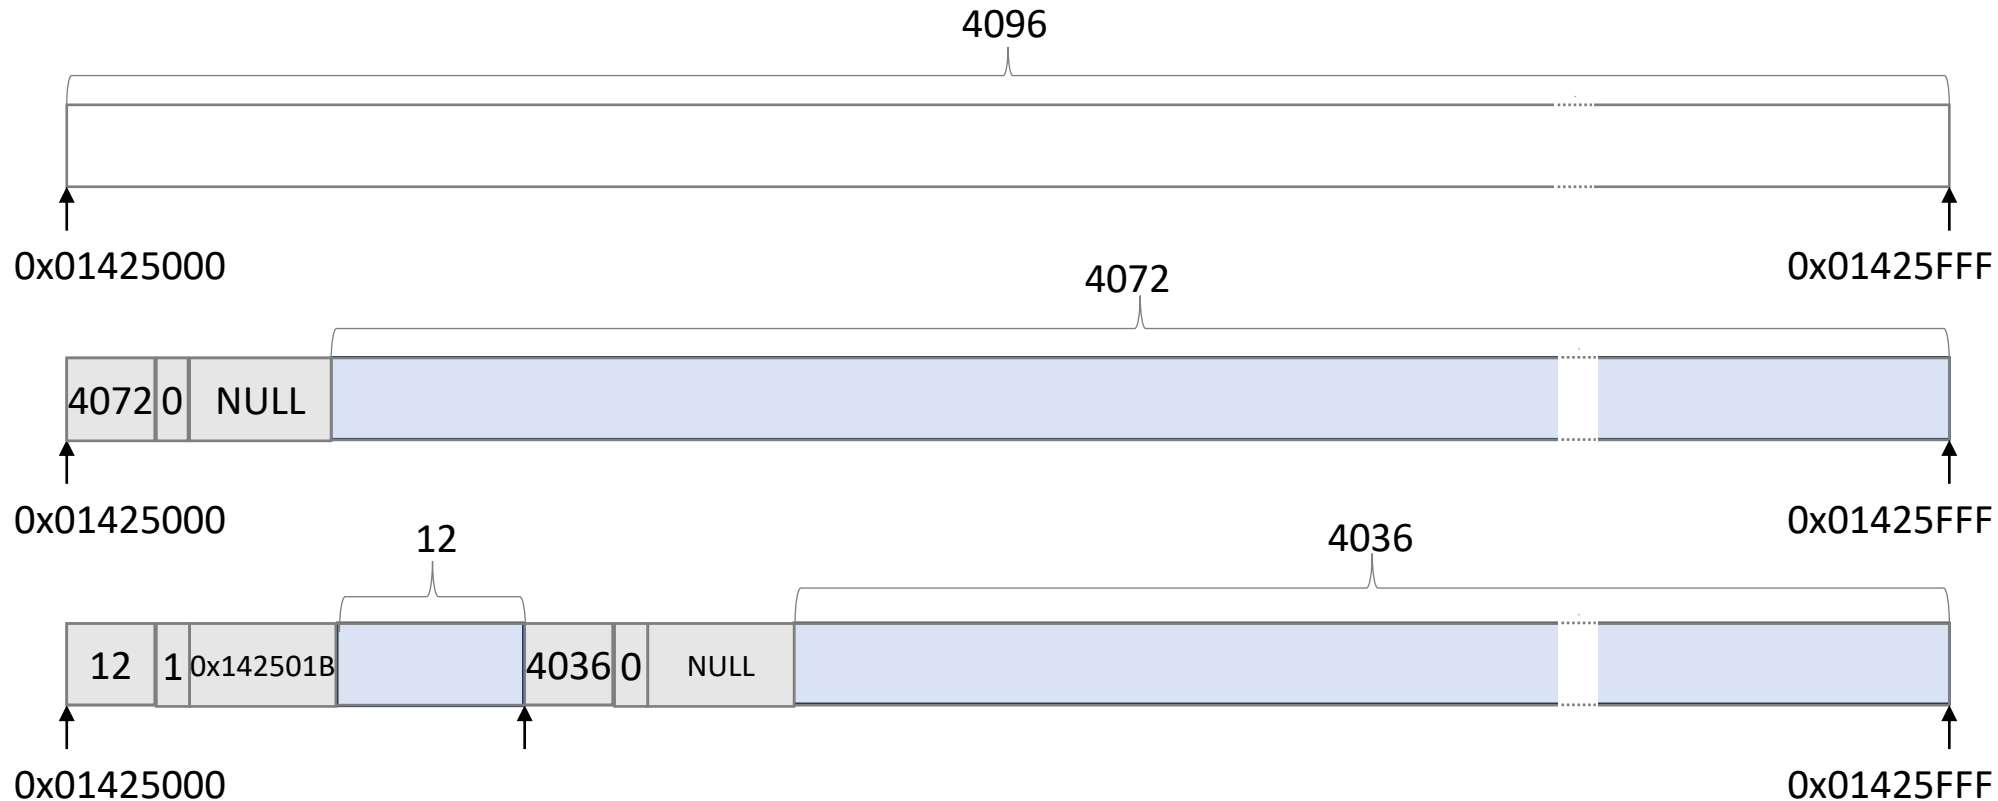
Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02

Memory Allocation Management (2/2)

- Data regions and headers are placed inside the pages acquired by `mmap()`
 - embedded linked list
- Pages by `mmap()` are fully filled with data regions and their headers



Homework 2.
malloc

5118020-03
Operating Systems

2024-05-02

Operations (1/2)

7

- `smalloc (size_t s)`
 - choose an unused data region whose size is greater than or equal to `s`
 - if no exists, use `mmap()` to add more pages
 - if the size of the chosen data region is greater than `s + 24`, split it into two
 - update the size of the chosen data region as `s`
 - add a new header for the remaining data region
 - return the first address of the data region
- `smalloc_mode (size_t s, smmode m)`
 - choose an unused data region with the specified mode `m`
 - mode: `bestfit`, `worstfit`, `firstfit`
 - the other behaviors are the same as `smalloc (size_t s)`

Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02

Operations (2/2)

8

- `srealloc (void * p, size_t s)`
 - find the header for the data region of pointer `p`
 - if it does not exist, abort the program execution (i.e., `abort()`)
 - resize the data region by `s`
 - change the data region if needed
 - split the data region if possible
- `sfree (void * p)`
 - find the header for the data region of pointer `p`
 - if it does not exist, abort the program execution (i.e., `abort()`)
 - mark the data region as unused at the header
- `smcoalesce ()`
 - if there exist two (or more) pairs of header and data region are unused and adjacent, merge them into one pair of header and data region

Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02

Submission

- All results must be submitted via LMS
 - Source code files
 - Submit all source code
 - You must provide a build script (e.g., bash script or Makefile) and its instruction document (e.g., README) if needed
 - Presentation
 - Submit the video record file; or, you can submit the URL to the presentation video on web
- No late submissions will be accepted

Video Presentation

- Take a 8-min video for reviewing the source code and testing the program
 - either in Korean or in English
 - every team member must take a part in presentation
- Your video must show how each API is implemented and used
 - devise scenarios for `smalloc_mode` and `smalloc_coalesce` and demonstrate their operations
 - suggest ideas to improve `smalloc` further

Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02

Notes

- Welcome your questions anytime on the Slack channel
- The team members must share the same responsibilities and take in charge of all tasks together
 - Peer evaluation follows immediately after the submission deadline
 - Inform me quickly if you keep fail to contact with your teammate
- It is strictly permitted to use auto-programming tools in any form

Homework 2.
smalloc

5118020-03
Operating Systems

2024-05-02