

5118020-03 Operating Systems

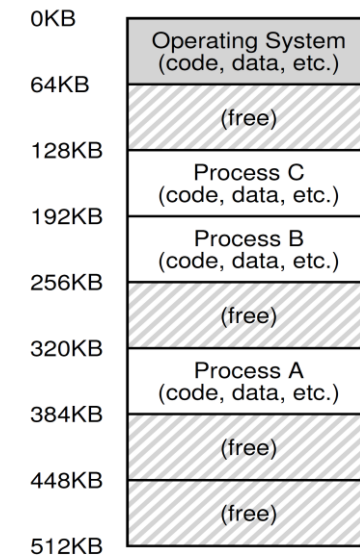
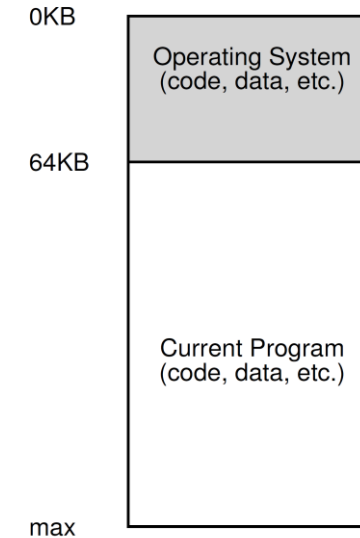
Address Space and Dynamic Relocation

OSTEP Chapters 13 & 15

Shin Hong

Motivation

- Early computer systems did not need memory abstraction since there was no issue for a program to occupy whole memory
- Memory abstraction is required with time-sharing
 - approach 1. like CPU context switching, store the entire memory state to a storage device at a context switching
 - heavy context switching cost
 - approach 2. let a process use only a region of memory, and keep multiple processes in the memory at the same time
 - low utilization of memory
 - data protection issue

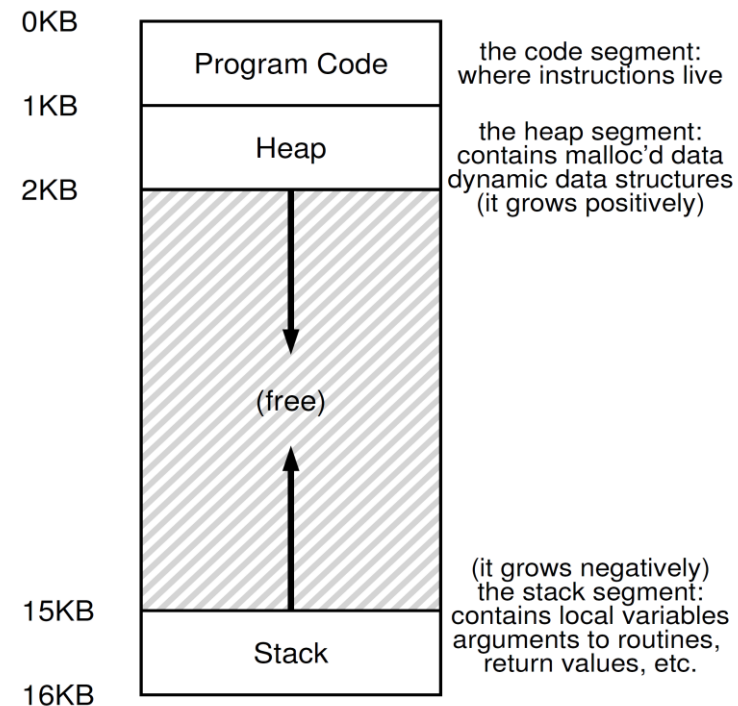


Address Space
and Dynamic
Relocation

Abstraction: The Address Space

3

- Address space is the running program's view of memory
 - interface between a process and memory devices
- The address space of a process has a continuous region of addresses which contains the code, the stack, the heap and all memory state



Address Space
and Dynamic
Relocation

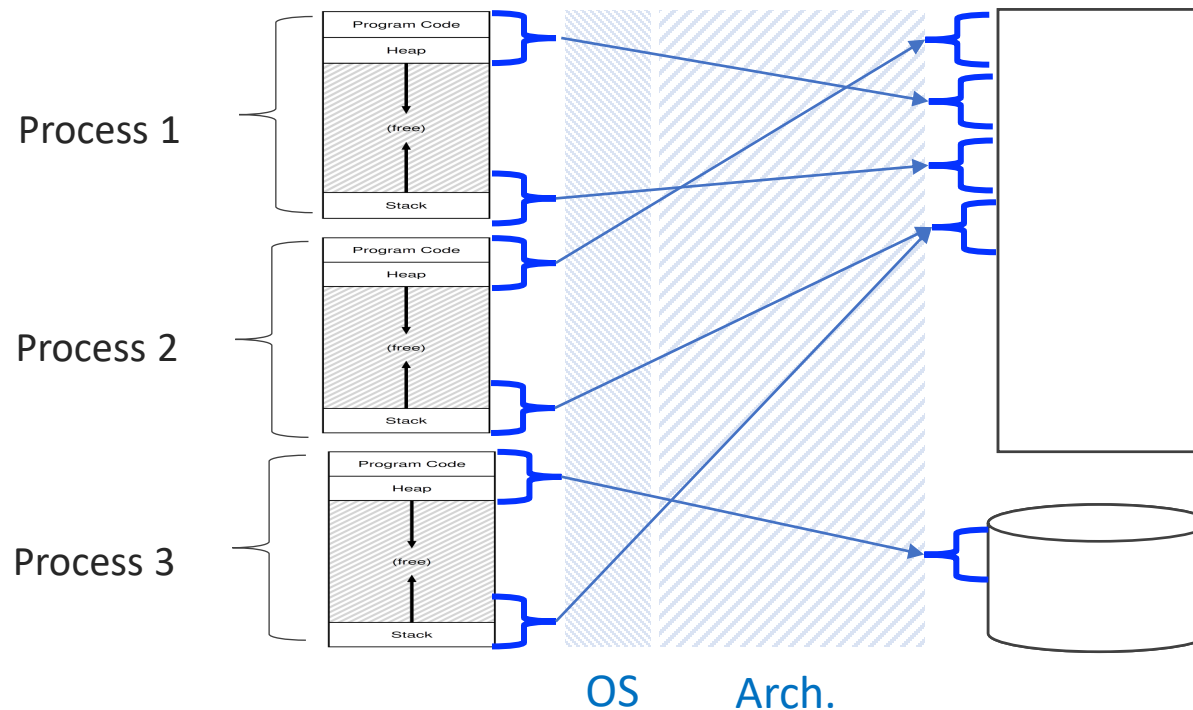
5118020-03
Operating Systems

2024-04-15

Virtual Memory

4

- The OS **virtualizes** memory in cooperation with computer architecture
 - the goals of memory virtualization
 - transparency (seamless-ness)
 - time-efficiency and space-efficiency
 - isolation



Address Space
and Dynamic
Relocation

5118020-03
Operating Systems

2024-04-15

Hardware-based Address Translation

5

- Let a computer architecture transform each memory access by converting a virtual address to a physical address
 - like a computer architecture translates relative addresses to absolute addresses
- The OS manages a mapping from virtual addresses to physical addresses
 - the OS interposes between an application program and hardware operation at critical points to maintain control over the hardware
 - the critical points include:
 - process creation/termination,
 - context switching,
 - when a process attempts to access forbidden memory regions

Address Space
and Dynamic
Relocation

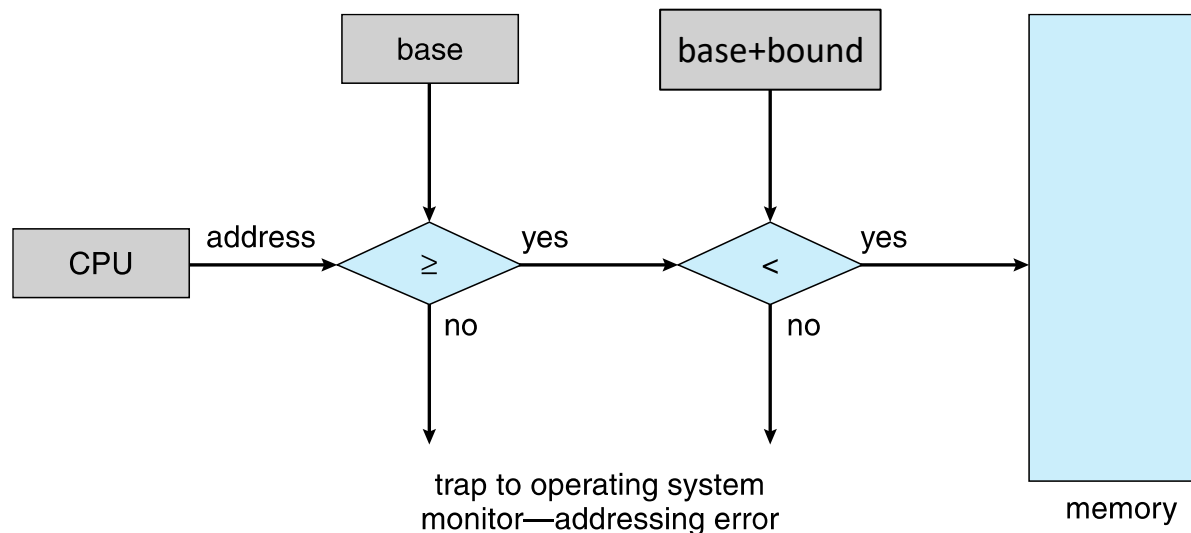
5118020-03
Operating Systems
2024-04-15

Approach 1: Dynamic Relocation

6

- Assumption

- The size of the address space for a process is much smaller than the total amount of available memory in the main memory device
- Every process is given the same amount of address space
- The MMU of the computer architecture supports the **base** register and the **bound** (limit) register
 - always translate a memory address if it's user mode
 - the base and the bound registers can be accessed only if it's in privileged mode



Address Space
and Dynamic
Relocation

5118020-03
Operating Systems
2024-04-15

Approach 1: Dynamic Relocation

7


- Approach

- Allocate a continuous region of physical memory to a process
- Store the beginning of the allocated memory region to the base register b
- Always translate a memory address of a program m into $m + b$
- Set the bound register to raise a trap if the process tries to access an address beyond its given capacity

```
void func() {  
    int x = 3000;  
    x = x + 3;  
    ...  
}
```

↓

```
128: movl 0x0(%ebx), %eax  
132: addl $0x03, %eax  
135: movl %eax, 0x0(%ebx)
```



Address Space
and Dynamic
Relocation

5118020-03
Operating Systems

2024-04-15

Cooperation of CA and OS

8

- Computer Architecture

- enforce address translation and bound check under user mode
- raise a trap at a bound violation
- disallow updating the base and the bound register under user mode

- OS

- split available physical memory into multiple memory slots
 - maintain a process table and a list of free memory slots
- allocate a free slot to a new process
- reclaim the used slot at a process termination
- update base at context switching
- handle a trap (exception) raised by bound check

Address Space
and Dynamic
Relocation

5118020-03
Operating Systems
2024-04-15

Example.

Limited Direct Execution & Dynamic Relocation

OS @ run (kernel mode)	Hardware	Program (user mode)
To start process A: allocate entry in process table alloc memory for process set base/bound registers return-from-trap (into A)	restore registers of A move to user mode jump to A's (initial) PC	Process A runs Fetch instruction
	translate virtual address perform fetch	Execute instruction
	if explicit load/store: ensure address is legal translate virtual address perform load/store	(A runs...)
	Timer interrupt move to kernel mode jump to handler	
Handle timer decide: stop A, run B call <code>switch()</code> routine save regs(A) to <code>proc-struct(A)</code> (including base/bounds) restore regs(B) from <code>proc-struct(B)</code> (including base/bounds) return-from-trap (into B)	restore registers of B move to user mode jump to B's PC	Process B runs Execute bad load
	Load is out-of-bounds; move to kernel mode jump to trap handler	
Handle the trap decide to kill process B deallocate B's memory free B's entry in process table		

9


Address Space and Dynamic Relocation

5118020-03
Operating Systems

2024-04-15

Limitations

- internal fragmentation
- the number of processes afforded in physical memory space
 - runtime cost of write-back at context-switching



10

Address Space
and Dynamic
Relocation