

5118020-03 Operating System

# Segmentation

OSTEP Chapters 16 & 17

Shin Hong


# Motivation

- **Internal fragmentation problem**

- a process typically uses a small portion of an address space in a sparse manner
- the rest of the address space remains unused and wasted

- **Redundant data**

- The same piece of code may be loaded redundantly if multiple processes use it



# Approach

3

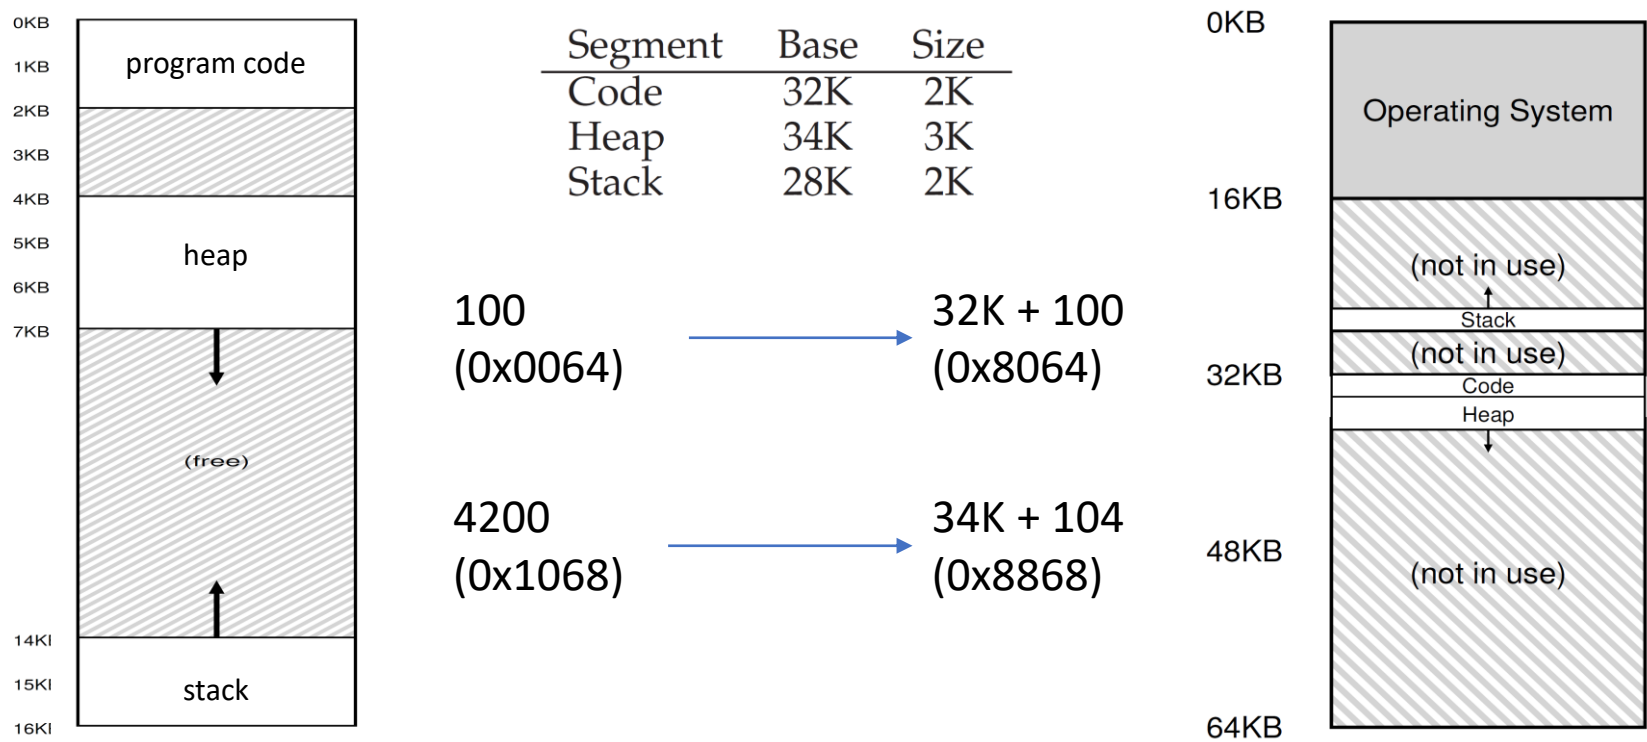
- Ideas
  1. split an address space into multiple pieces and manage each separately
  2. allocate the memory to each piece depending on the actual needs
- Define the address space of a process as a set of **segments**
  - a segment is a continuous memory region defined by a pair of base and bounds addresses
  - the available portion of the address space of a process can be represented as a set of segments
  - a segment is initially defined at a loading time
  - a segment can be relocated, extended, or shrunk over time
- Possible to make multiple processes share the same segment if they use the same code or data

Segmentation

5118020-03  
Operating System  
2024-04-15

# Segmentation

- A process has three segments, **code**, **heap** and **stack** in an address space
- MMU holds three pairs of base-bound values, and for each memory access, it identifies which segment the access is on
  - For a virtual address, MMU finds the base value, and add it to the offset part to obtain the corresponding physical address
- On demand, OS can change bounds to extend/shrink segments of a process




# Which Segment A Memory Access is On?

5

- By virtual address

- use few most significant bits as a segmentation indicator
- ex. virtual address 4200



- By instruction type

- referring the code segment if the address is derived from PC
- referring the stack segment if the address is derived from stack pointer
- referring the heap segment otherwise

Segmentation

5118020-03  
Operating System

2024-04-15

# Segment Attributes

6

- A bit to indicate whether a segment grows forward or backward

- ex.

Segment	Base	Size (max 4K)	Grows Positive?
Code <sub>00</sub>	32K	2K	1
Heap <sub>01</sub>	34K	3K	1
Stack <sub>11</sub>	28K	2K	0

- A bit to indicate whether a segment is for read-write or read-only

- ex.

Segment	Base	Size (max 4K)	Grows Positive?	Protection
Code <sub>00</sub>	32K	2K	1	Read-Execute
Heap <sub>01</sub>	34K	3K	1	Read-Write
Stack <sub>11</sub>	28K	2K	0	Read-Write

Segmentation


5118020-03  
Operating System

2024-04-15

# Operating System Supports

7

- context switching
- serve requests of growing/shrinking a segment
- manage free space to counter external fragmentation
  - compaction
  - free-list management



Segmentation


5118020-03  
Operating System

2024-04-15

# Free-Space Management

8

- External fragmentation hurts memory utilization
  - commonly happens for dynamic allocation of variable-length memory units
  - example. with a 30-bytes address space
    - alloc 15 bytes as A
    - alloc 10 bytes as B
    - free A (15 bytes)
    - alloc 5 bytes as C
    - alloc 15 bytes as D



Segmentation

5118020-03  
Operating System

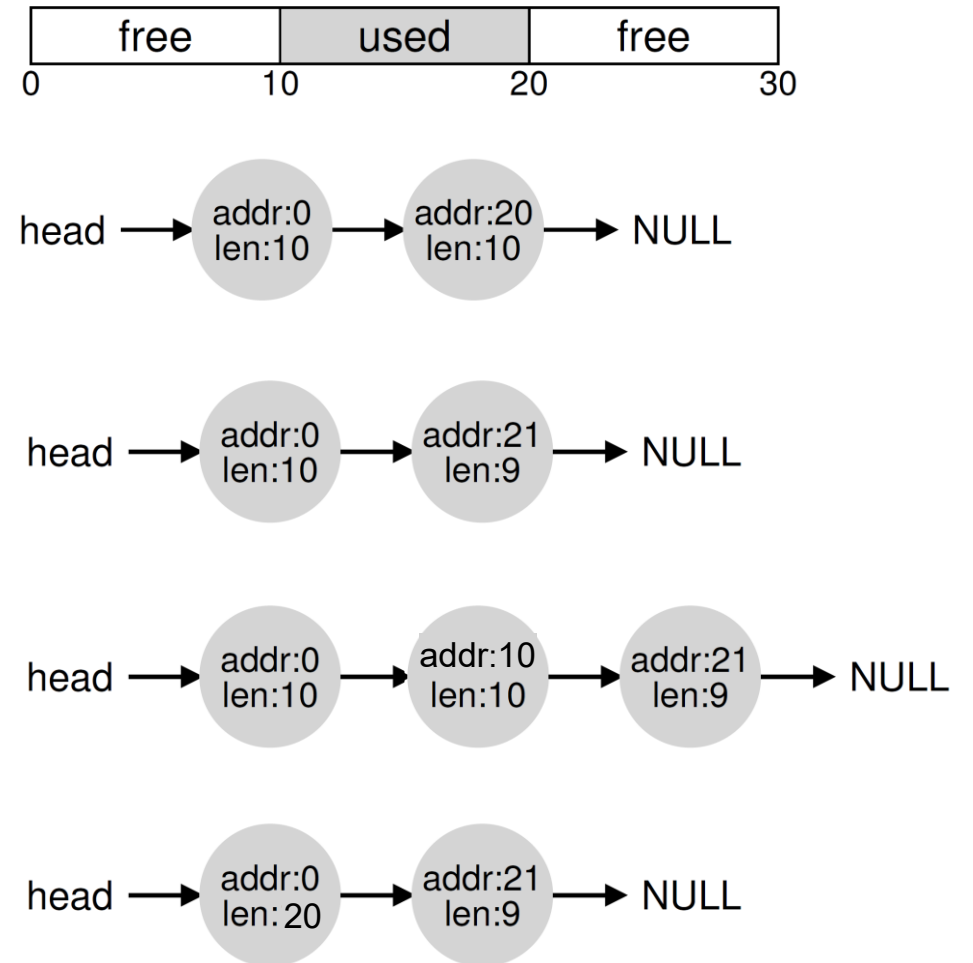
2024-04-15



# Free List

9

- maintains free space as a linked list of free chunks (i.e., continuous unused memory regions)
- to allocate a chunk of size  $m$ 
  - find a node of a free chunk whose size is greater than or equal to  $m$
  - split the free chunk if its size is greater than  $m$
- to free an allocated memory chunk,
  - add a node of the newly freed chunk to the free list
  - merge adjacent chunks into a single larger chunk (coalescing)




# Allocation Strategies

10

- The ideal allocator should be fast and minimize fragmentation.
- **Best fit:** search through the free list and find the smallest chunks that are large enough to afford the memory request (i.e., closest to what the user asks)
- **Worst fit:** find the largest chunk in order to keep large chunks remain in the free list
- **First fit:** find the first chunk that is large enough to afford the requested memory
- **Next fit:** conduct the first fit search from the node where the last search was stopped (not from the beginning of the free list)

Ex. the user asks 18



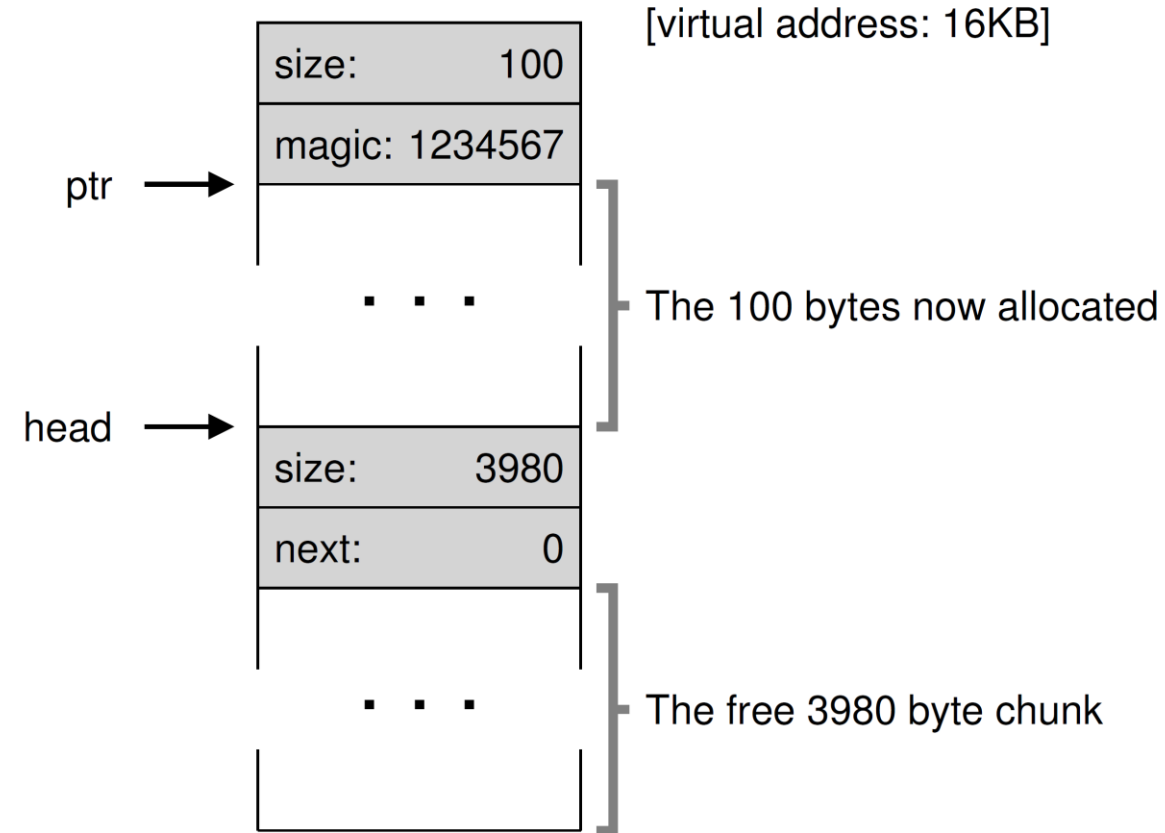
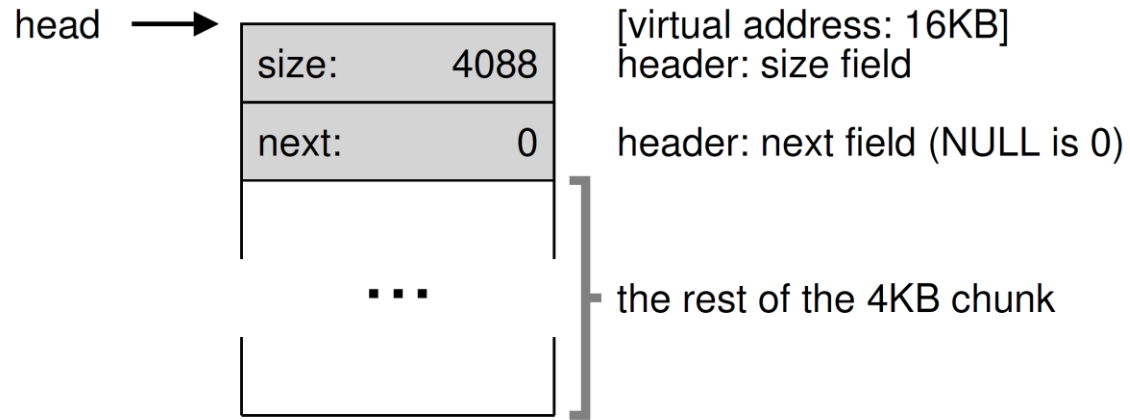
Segmentation

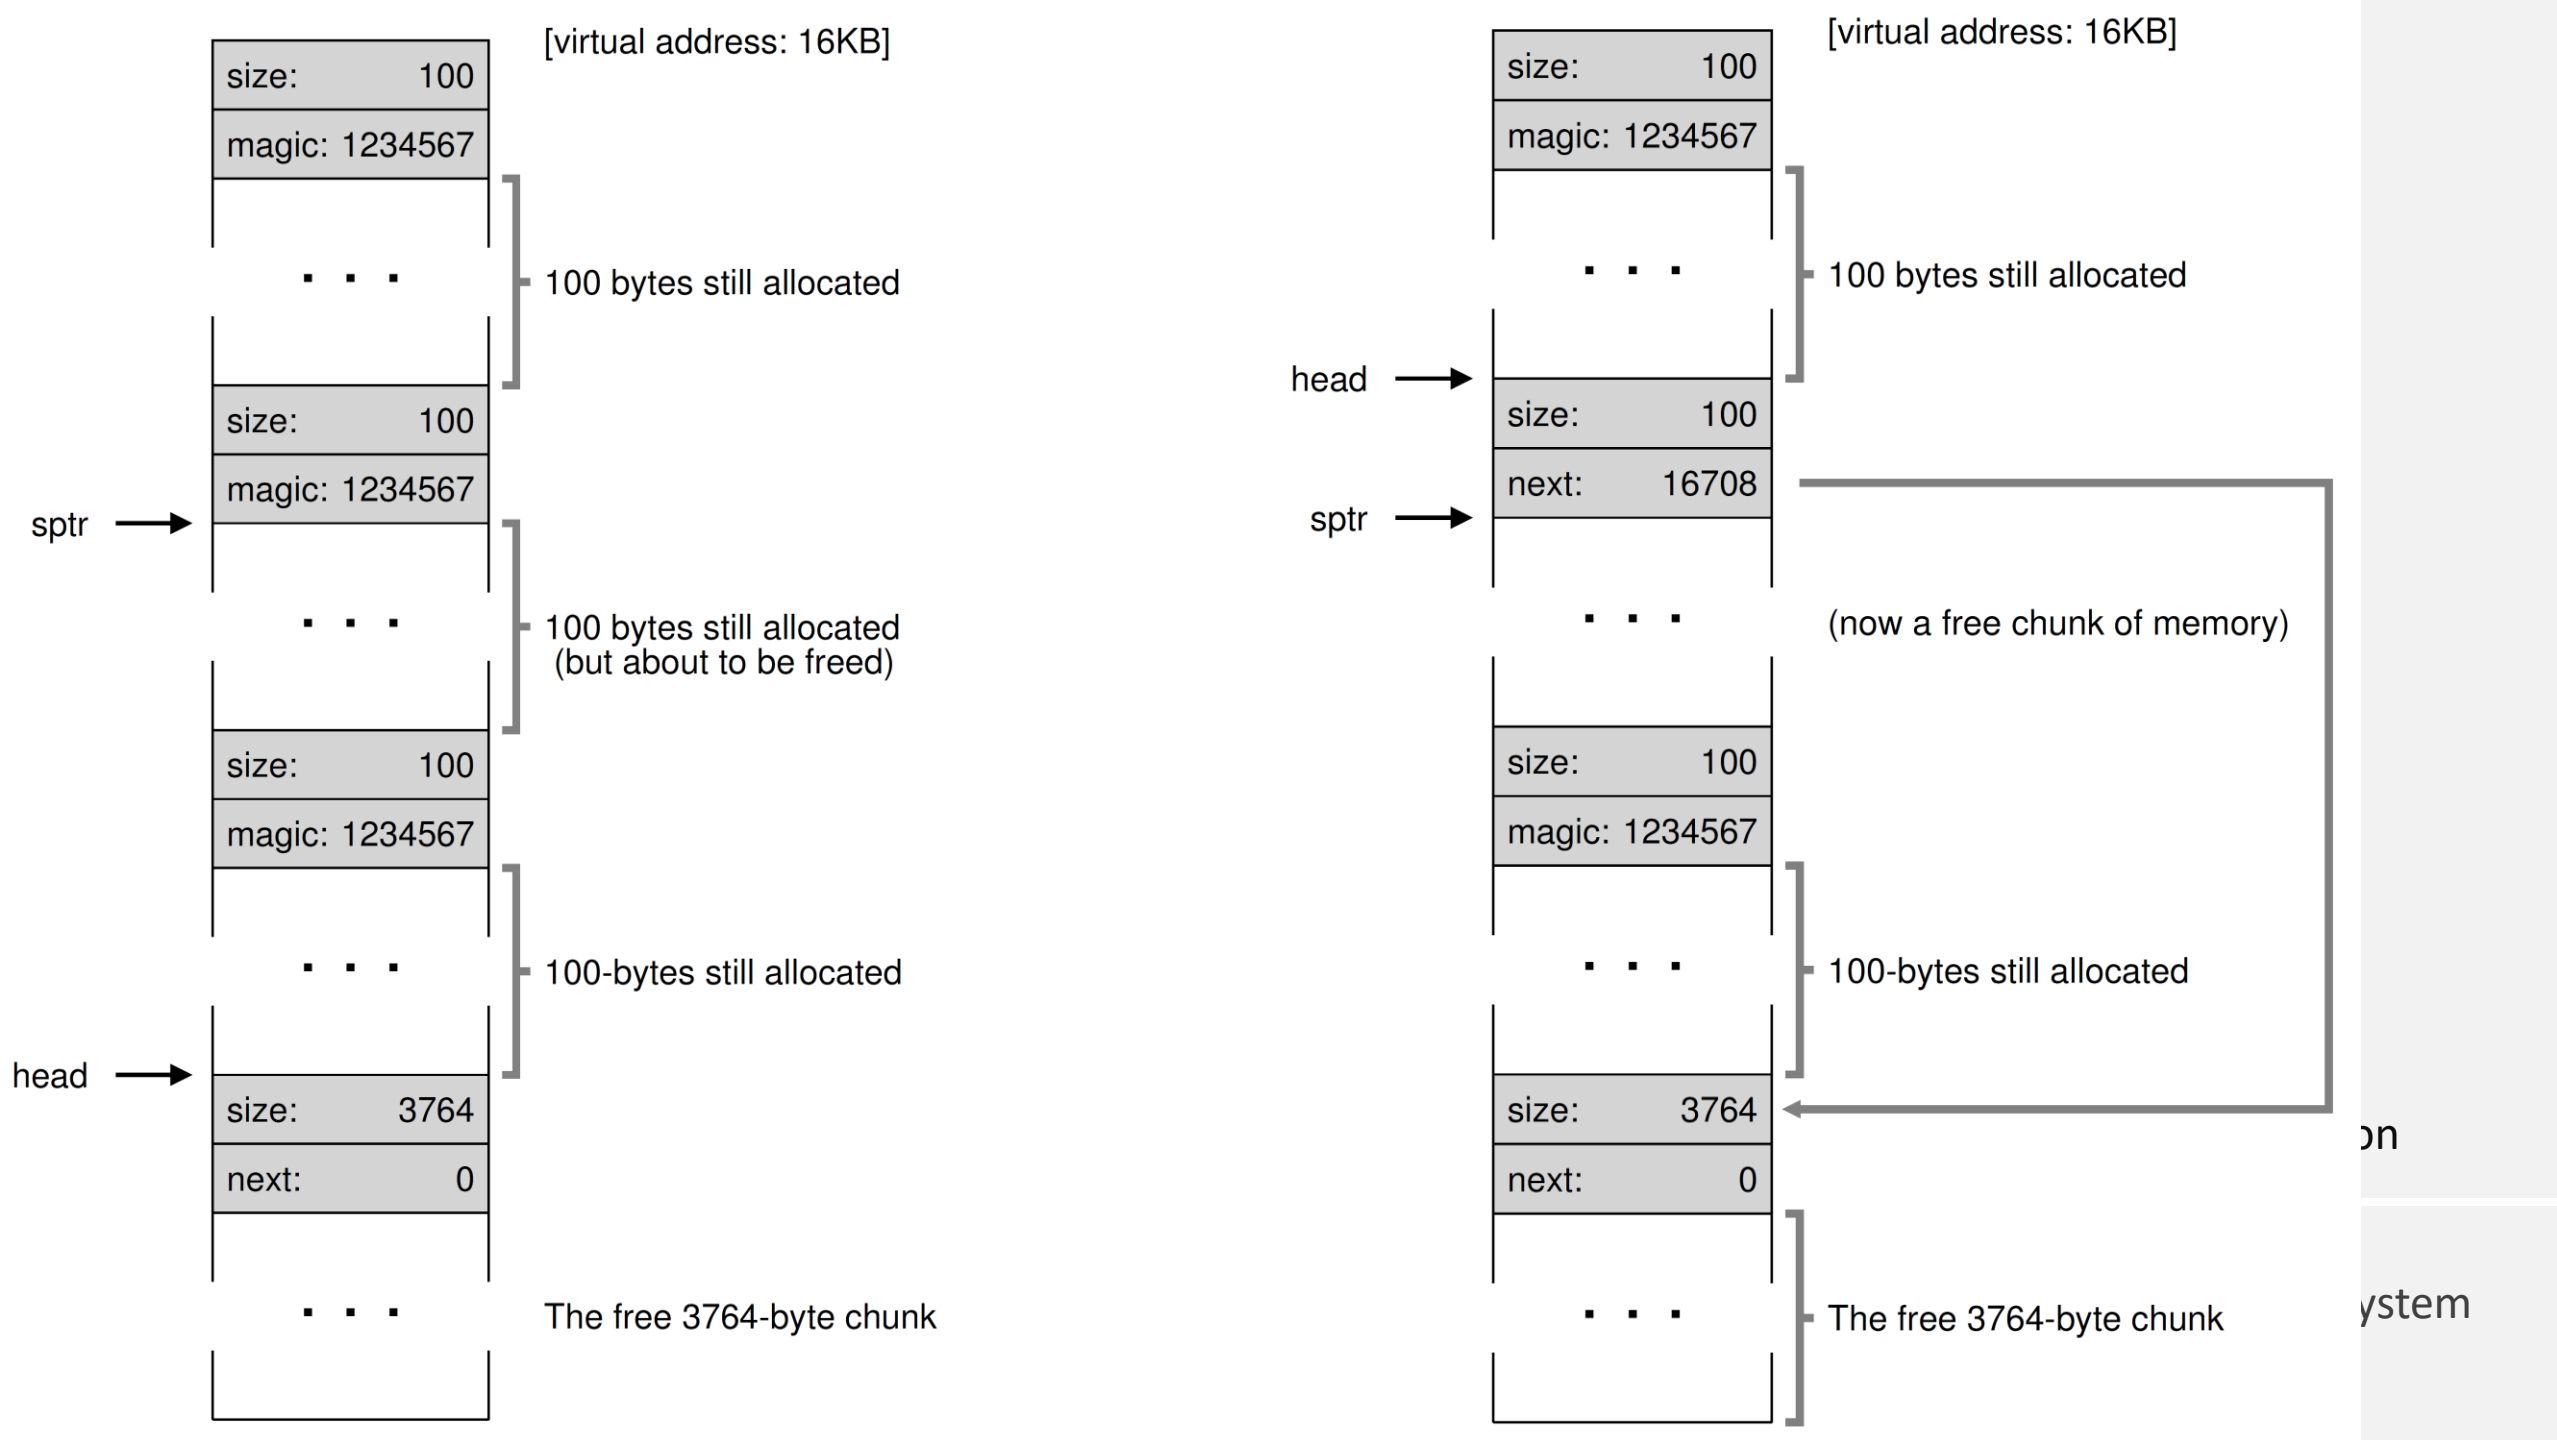
5118020-03  
Operating System

2024-04-15

# Embedding A Free List

- Example: 4096-byte memory





# Buddy System (1/3)

- Use **power-of-2 allocator**
  - satisfies requests in units sized as a power of 2 (e.g., 16, 32, 64)
  - request rounded up to next smallest power of 2 (e.g., 32 for 20, 64 for 40)
  - split an existing chunk into two buddies when the smaller chunk is needed
    - continue until an appropriately-sized chunk is available
- E.g., assume one chunk of 256KB is available, there is a request of 21 KB
  - split the chunk into  $A_L$  and  $A_R$  of 128 KB each
  - split  $A_L$  into  $B_L$  and  $B_R$ , each of which has 64 KB
  - split  $B_L$  into  $C_L$  and  $C_R$  of 32 KB each – one used to satisfy request

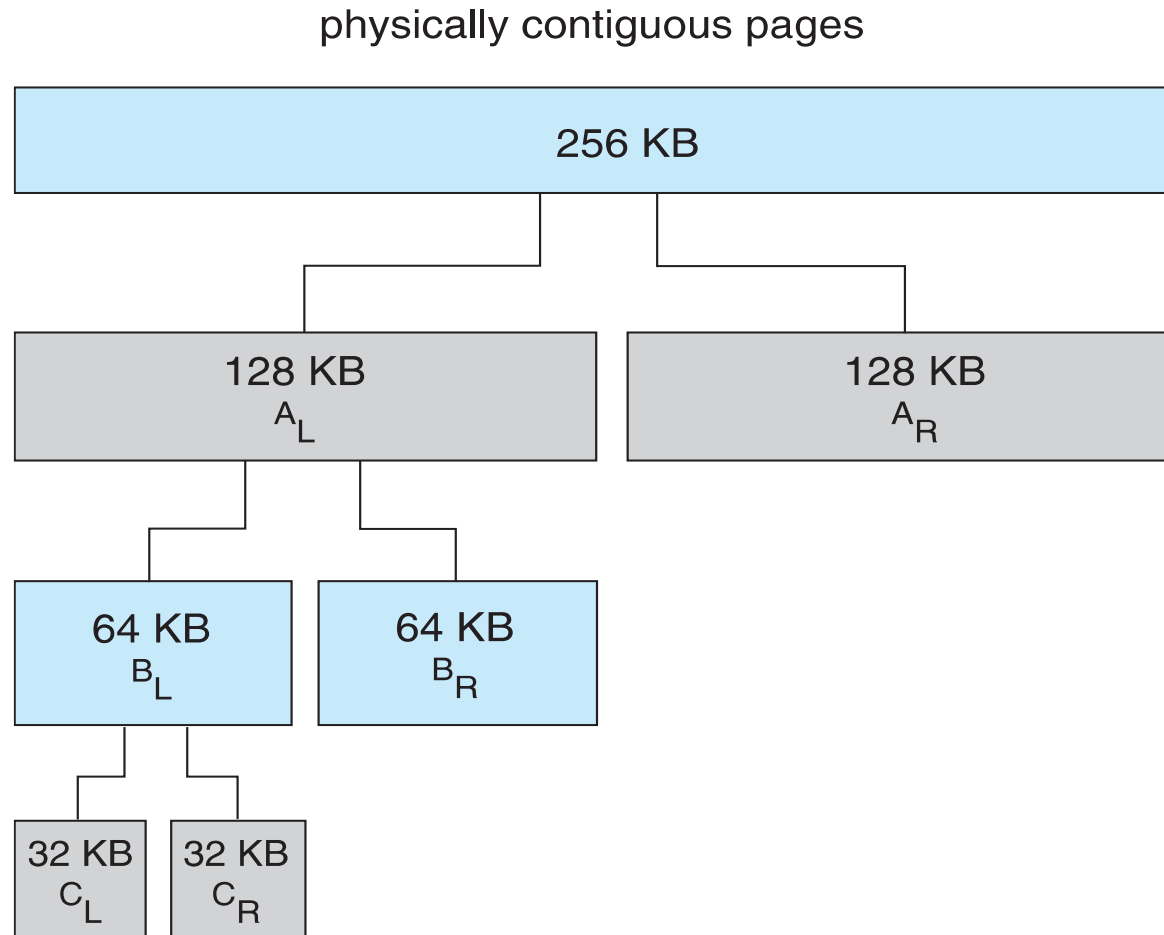
Segmentation

5118020-03  
Operating System

2024-04-15

# Buddy System (2/3)

14



Segmentation

5118020-03  
Operating System

2024-04-15

# Buddy System (3/3)

15

- Advantage: quickly **coalesce** unused chunks into larger chunk
- Disadvantage: internal fragmentation

Segmentation

5118020-03  
Operating System

2024-04-15