

5118020-03 Operating System

Paging

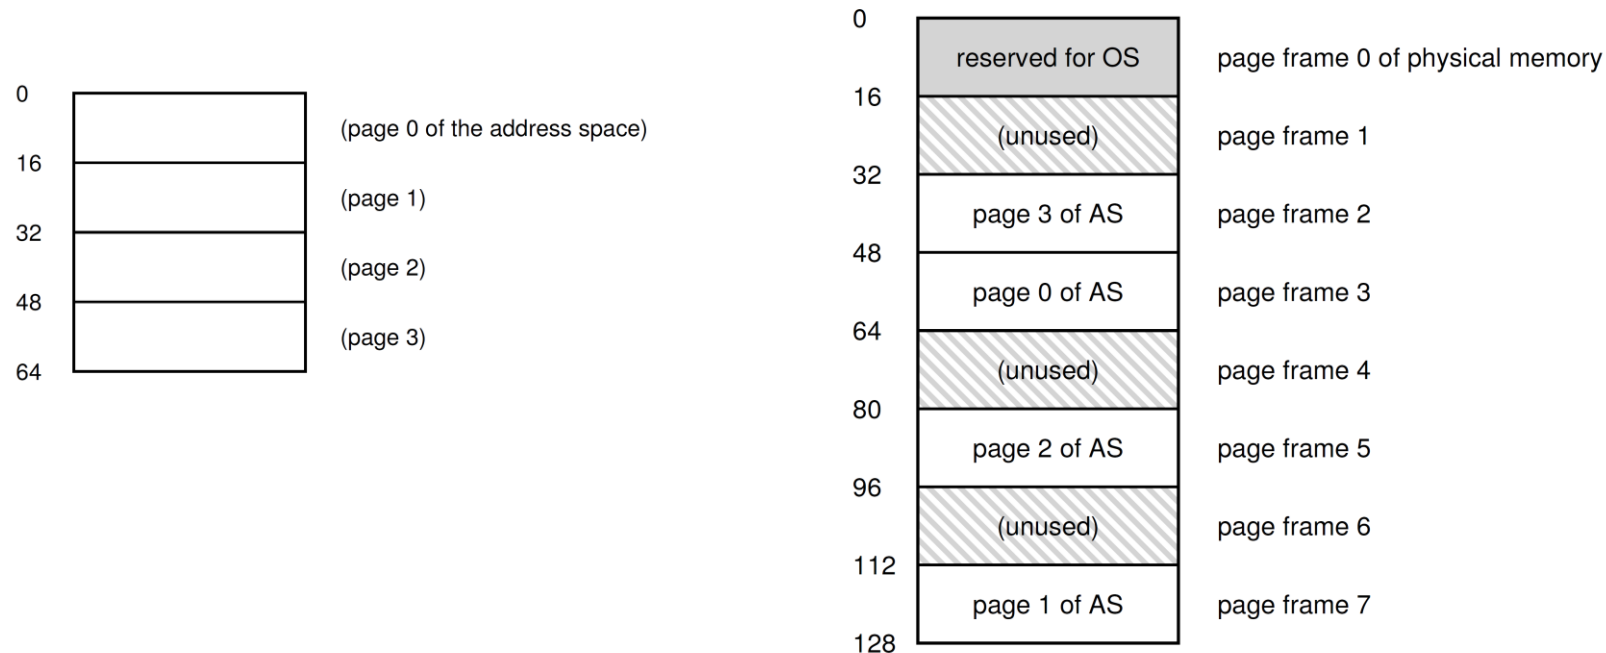
OSTEP Chapters 18 and 19

Shin Hong

Paging

2

- Paging is to allocate memory to a process as a set of small fixed-size pieces
 - Divide an address space into **pages** and a physical memory space into **frames** such that a page and a frame have the same size
 - Limit both internal and external fragmentation
 - it is empirically found that first-fit loses 33% of memory due to external fragmentation
 - Ex. 16-byte pages with 64 bytes address space and 128 bytes physical memory



Paging

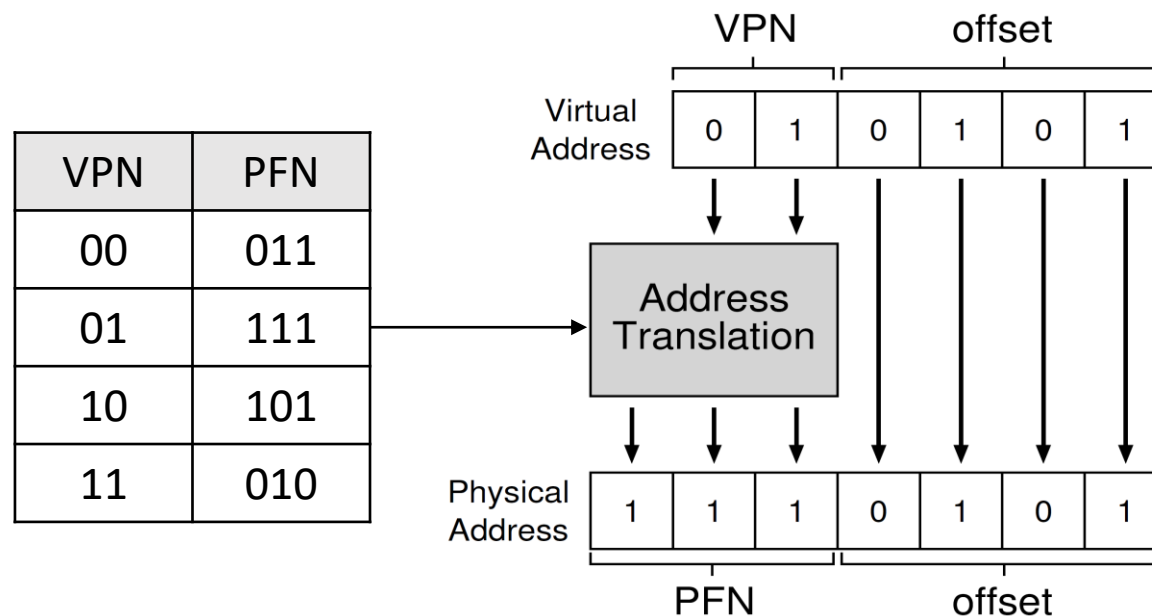
5118020-03
Operating System

2024-04-23

Page Table and Address Translation

3

- A page table maintains to which frame a page is assigned
 - as per-process page table, or as inverted page table
- In an address translation, a virtual address is split it into a VPN and an offset, and then the VPN is replaced with the corresponding PFN
 - ex. 16-byte page, 64-byte address space & 128-byte physical memory



Paging


5118020-03
Operating System

2024-04-23

Per-process Page Table Entry

4

- Physical Frame Number (PFN)
 - valid bit: whether the page is given to the process (i.e., in use)
 - present bit: whether the page is in physical memory or on disk
 - permission bits: read-only, read-write, super mode-only, etc.
 - dirty bit: whether the page was modified since it was brought into memory
 - reference bit: whether the page has been recently accessed
-
- Example of 32-bit x86



Paging

5118020-03
Operating System


2024-04-23

Accessing Memory with Paging

5

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

VPN	PFN	
00	011	
01	111	
10	101	
11	010	



Paging

5118020-03
Operating System

2024-04-23


Ex. Memory Trace

6

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

↓

```
1024 movl $0x0, (%edi,%eax,4)  
1028 incl %eax  
1032 cmpl $0x03e8,%eax  
1036 jne 0x1024
```



Page Table Size

7

- page tables take a large amount of memory
 - example
 - 4 KB page/frame in a 32-bit address space and a 4 GB physical mem.
 - 20-bits VPN, thus 2^{20} entries in a page table
 - assume that 4 bytes is needed for each entry, thus, 4 MB ($=2^{22}$ bytes) for a page table
 - 400 MB for 100 processes
- a page table should be resided in a main memory since MMU cannot accommodate a whole page table at a time
 - every memory access incurs at least one extra memory access for address translation

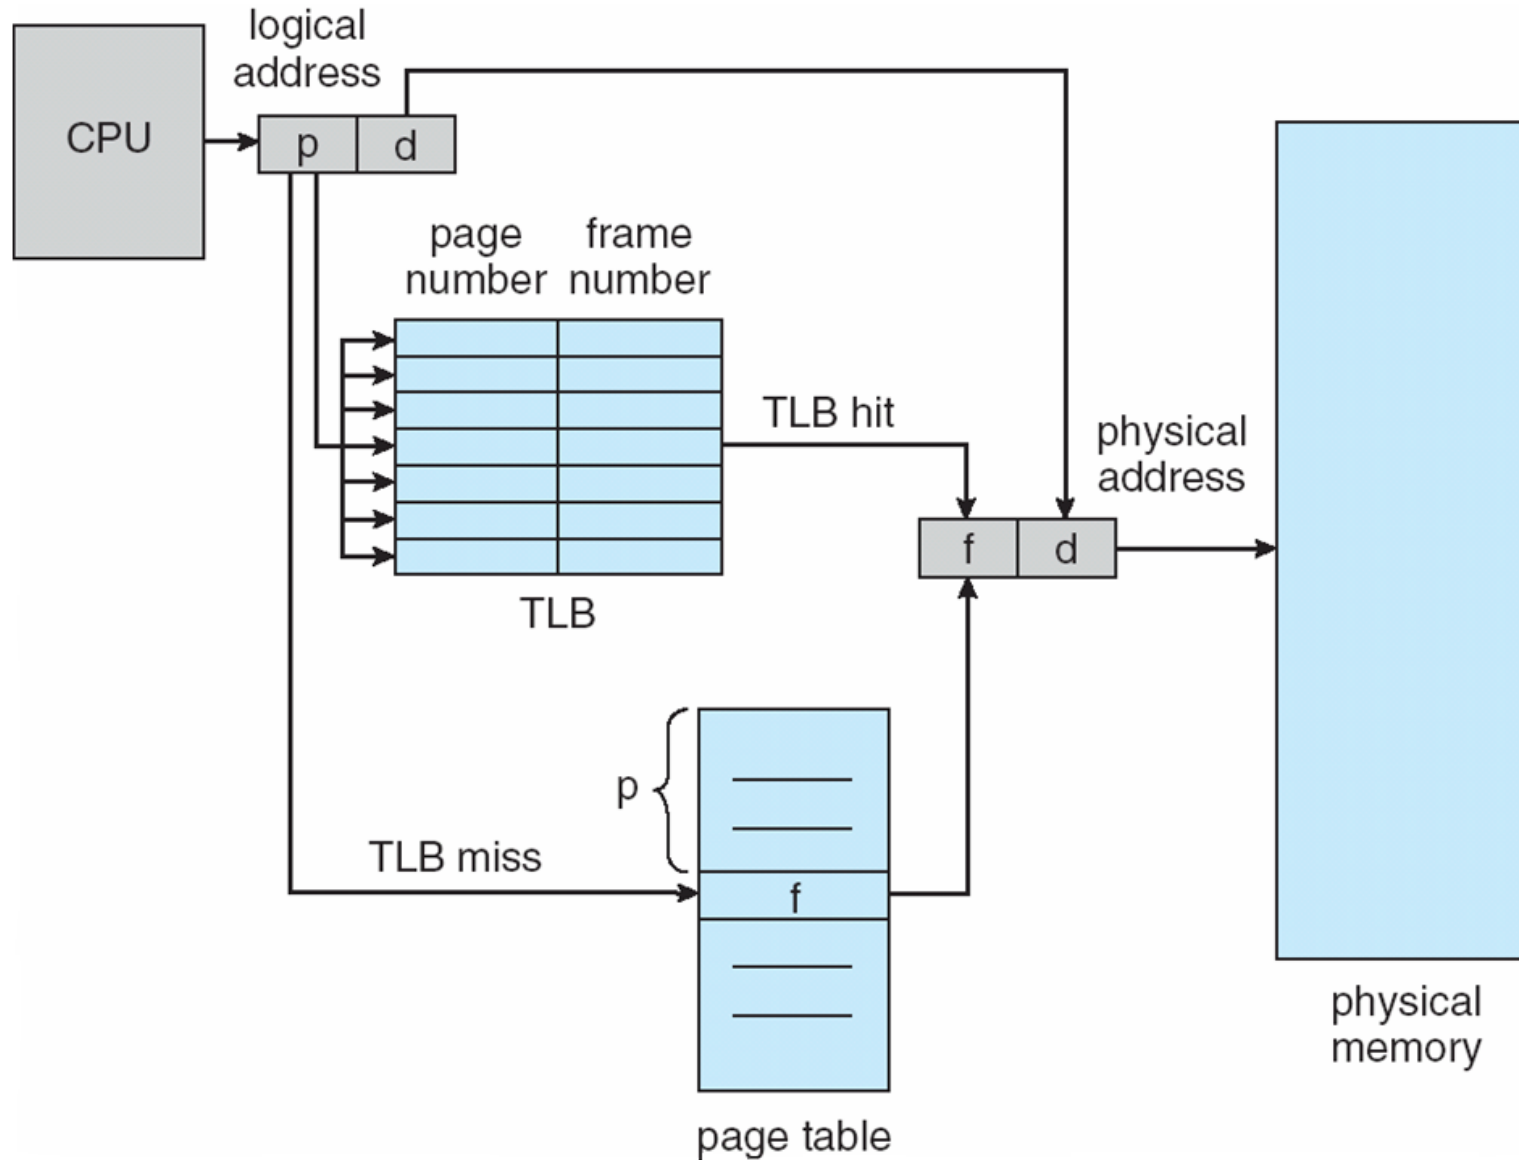
Paging

5118020-03
Operating System

2024-04-23

Paging Hardware with TLB

8



Paging

5118020-03
Operating System

2024-04-23

Translation-lookaside Buffer (TLB)


9

- A TLB is a part of MMU working as a cache of a page table
 - upon a memory access request, the computer architecture first checks the TLB to see if it has page translation information for the corresponding VPN
 - address-translation cache
- Steps for translating a virtual address with TLB
 - extract the VPN from a virtual address
 - check if the entry for the VPN is found in the TLB
 - if there exists, get the PFN from the TLB (i.e., TLB hit)
 - otherwise (i.e., TLB miss)
 - reference the page table to get the PFN
 - update the TLB with the VPN and the PFN
 - repeat from the beginning

Paging

5118020-03
Operating System
2024-04-23

Memory Reference Pattern



11

Paging

5118020-03
Operating System

2024-04-23

Locality

12

- A TLB can save memory accesses because an application program tends to make memory accesses with **temporal locality** and **spatial locality**
 - **temporal locality**: if a program accessed a memory at address x , it will likely access x again in near future.
 - **spatial locality**: if a program accesses a memory at address x , it will likely access $x + d$ in near future (where d is a small number)
- Trade off between TLB size and TLB access speed
 - The bigger the size of a TLB is, the higher the hit ratio is
 - The bigger the size of a TLB is, the longer the look-up time takes

Paging

5118020-03
Operating System

2024-04-23

Handling TLB Miss

13

- Hardware-managed TLB
 - there must be a special register to point to the page table
 - the structure of a page table must be fixed
- Software-managed TLB
 - the architecture raises an exception in a middle of an instruction execution if a TLB miss occurs
 - the OS finds proper information from page tables and updates TLB
 - after the trap handler execution, the architecture retries the instruction

Paging

5118020-03
Operating System

2024-04-23

TLB Entry

14

- A TLB of an architecture typically has 32, 64 or 128 entries
 - parallel search
- A TLB entry typically has the following attributes
 - VPN
 - PFN
 - valid bit
 - protection bit: read-only or read-write
 - dirty bit

Paging

5118020-03
Operating System

2024-04-23

Managing TLB at Context Switching

15

- The existing TLB entries must not be used after a context switching
- Approaches
 1. Flush all entries
 - flushing can be made by turning off the valid bit of every entry
 - a series of TLB misses will happen right after a context switch
 2. Have a process identifier in a TLB entry
 - address space identifier (ASID)
 - multiple processes can readily share a TLB

Paging

5118020-03
Operating System

2024-04-23

TLB Entry Replacement Policies

16

- Which entry to evict from a full TLB when a new entry must be installed?
 - decision should be made to minimize a TLB miss rate
 - a typical case of the cache replacement problem
- Approaches
 1. evict the least-recently-used (LRU) one
 2. evict a random entry

Paging

5118020-03
Operating System

2024-04-23

Least Recently Used (LRU) Algorithm

17

- Use past knowledge to predict future
- Associate time of last use with each page
- Choose one that has not been used in the most amount of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0		1		1		1
	0	0	0		0		0	0	3	3		3		0		0
		1	1		3		3	2	2	2		2		2		7

page frames

Paging

5118020-03
Operating System

2024-04-23