

5118020-03 Operating System

Swapping

OSTEP Chapters 21 and 22

Shin Hong

Beyond Physical Memory

2

- What if the amount of the allocated pages in running processes exceed the physical memory capacity?
- OS stashes away some pages that are not in great demand at the moment
 - usually to a swap space in HDD or SSD
- OS brings a stored pages back to the main memory before a process accesses data on the page
 - by stashing other pages

Swapping

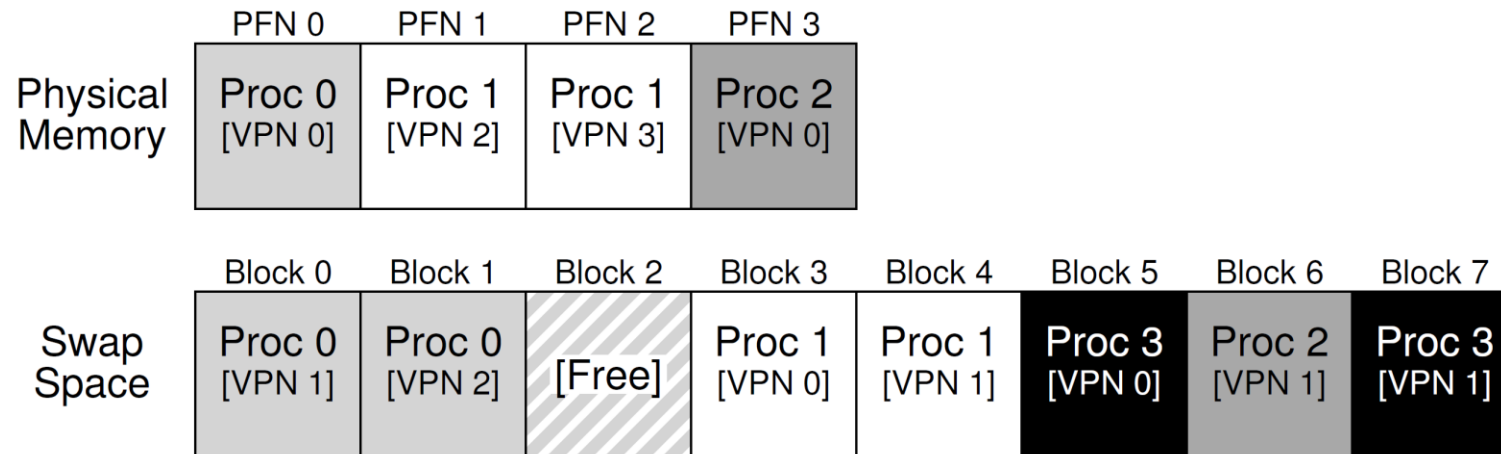
5118020-03
Operating System

2024-05-31

Swap Space

3

- reserve **swap space** on the disk for moving pages back and forth
 - swap space is a collection of page-sized blocks
 - OS needs to store the location on a swap space where a page is stored
- Ex.



Swapping

5118020-03
Operating System

2024-05-31

Page Fault

4

- The present bit of a page table entry indicates whether the page is present in physical memory, or stored in swap space
- A page fault is raised by an architecture if the present bit is off, so that the page-fault handler is invoked to serve the page fault
 - the location of the page in swap space is stored at a page table
 - the process will be blocked during the I/O for stashing and reloading
- Page fault and consequent page-in and page-out operations are all taken placed transparently to the process

Swapping

5118020-03
Operating System

2024-05-31

Page Fault Control Flow Algorithm

5

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else    // TLB Miss
11     PTEAddr = PTBR + (VPN * sizeof(PTE))
12     PTE = AccessMemory(PTEAddr)
13     if (PTE.Valid == False)
14         RaiseException(SEGMENTATION_FAULT)
15     else
16         if (CanAccess(PTE.ProtectBits) == False)
17             RaiseException(PROTECTION_FAULT)
18         else if (PTE.Present == True)
19             // assuming hardware-managed TLB
20             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
21             RetryInstruction()
22         else if (PTE.Present == False)
23             RaiseException(PAGE_FAULT)
```

Swapping

5118020-03
Operating System

2024-05-31

Page Replacement

6

- Before page in a page, OS must page out some pages to make a room if the memory is almost fully used
 - most OS's have high and low watermarks for a page daemon to proactively start page out in background
- OS must have a proper page replacement policy because unnecessary page-in or page-out incurs significant runtime cost
 - A disk access takes 10000 to 100000 times longer than a memory access
 - Average Memory Access Time ($AMAT$) = $T_{Mem} + p_{Miss} \times T_{Disk}$
 - E.g., $T_{Mem}=100$ ns, $T_{Disk}=10$ ms, $p_{Miss}=10\%$: $AMAT = 1.01$ ms

Swapping

5118020-03
Operating System

2024-05-31

Page Replacement Policies

7

- Policies

- FIFO
- Random
- Optimal (ideal)
- Least Recently Used (LRU)
- Least Frequently Used (LFU)

- Simple cost model

- There's only one process in the system, which uses all frames
- A workload is represented as a sequence of VPN accesses
- A page-out occurs only when the memory is full
- The time cost of workload is measured as the total number of page-in and page-out operations with the workload

Swapping


5118020-03
Operating System

2024-05-31

FIFO

8

- Page out the one that was paged in first (i.e., stayed longest time)
- Example: 3 frames, 4 pages (0 to 3)




Swapping

Random

9

- Pick a random page at a page replacement decision
- Example: 3 frames, 4 pages (0 to 3)




Swapping

Optimal

10

- Select a page whose next use is furthest in future
- Example: 3 frames, 4 pages (0 to 3)




Swapping

Least Frequently Used (LFU)

11

- Select a page which has been used least
- Example: 3 frames, 4 pages (0 to 3)




...apping

Least Recently Used (LRU)

12

- Select a page whose last use is furthest in past
- Example: 3 frames, 4 pages (0 to 3)




Swapping

Workload Example – No Locality

13

- A random workload with 10000 accesses on 100 pages



Swapping


5118020-03
Operating System

2024-05-31

Workload Example – 80:20 Workload

14

- 10000 accesses on 100 pages where 80% of accesses are on a 20% of the pages



Swapping


5118020-03
Operating System

2024-05-31

Workload Example – Looping Sequential

15

- 10000 accesses on 50 pages
- repeating sequential accesses on 0, 1, 2, ..., 49 for 200 times



Swapping

5118020-03
Operating System

2024-05-31

Approximating LRU

16

- An exact implementation of LRU is expensive
 - update last access time of a page at every memory access
- Clock (second-chance) algorithm using reference bit
 - a reference bit (use bit) is given to each page table entry, such that the reference bit is set to 1 by the hardware at a memory access
 - pages are maintained in a circular list with one clock hand
 - the clock hand iterates over the list to choose the first node with the reference bit off
 - the clock hand turns off the reference bit when it passes over the node with the reference bit on.



Swapping

5118020-03
Operating System

2024-05-31

Clock Algorithm Example

17




Swapping

5118020-03
Operating System

2024-05-31

Clock Algorithm Performance

18



Swapping

5118020-03
Operating System

2024-05-31

Considering Modified Bit

19

- A page table entry has a modified bit (dirty bit) which is turned on by a hardware if the page is ever written
- It is more efficient to evict a clean page than a modified page, because there is no need to write back (page out) the content to the storage
 - (reference bit, modified bit)
 - (0, 0) : best page to replace
 - (0, 1) : not quite as good, must write out before replacement
 - (1, 0) : probably will be used again soon
 - (1, 1) : probably will be used again soon and need to write out

Swapping

5118020-03
Operating System

2024-05-31

Other Virtual Memory Policies for Efficiency

20

- **Demand paging:** bring the page into memory when it is to get accessed
- **Prefetching:** bring the page ahead of its use time if it is likely to be used soon
- **Clustering:** collect multiple dirty pages and write them once to the storage, rather than writing each one by one

Swapping


5118020-03
Operating System

2024-05-31

Thrashing

21

- A system is in **thrashing** when it is trapped by serial pagings for the working sets of the processes exceed an available number of frames
 - page out an existing page, but quickly the page gets paged in again
 - frequent swapping results low CPU utilization, which leads the system to increase the degree of multiprogramming
- A system resolves thrashing by running only a subset of the processes or killing some of the running processes



Swapping

5118020-03
Operating System

2024-05-31